

How ChatGPT, Perplexity, and Gemini Decide Which Businesses to Recommend

April 24, 2026 · 8 min read · Hey Pearl Agency LLC

When an AI recommends a business, it isn't guessing. It's weighing signals. Understand the signals and you can influence the outcome.

Ask ChatGPT to recommend a bookkeeping firm, and it will name two or three. It will not name yours unless something specific has happened to make it confident that yours belongs on the list. That confidence is not random. It is the output of a set of signals the engine can measure, and those signals are learnable.

This is the part most operators miss. AI recommendations feel like a black box, so people assume there's nothing to be done. There is a great deal to be done. You just have to understand what the machine is actually weighing.

The Core Recommendation Signals

Across the major engines, the same handful of signals do the work. The weighting differs by platform, but the ingredients are consistent.

Entity Authority

Before an engine can recommend you, it has to be certain you exist as a distinct thing. Entity authority is the degree to which an AI recognizes your business as a real, well-defined entity with a stable identity across the web. Businesses that live in the knowledge graph — with a Google Knowledge Panel, a Wikidata entry, consistent profiles — clear this bar easily. Businesses that are just a website with a logo often don't. If the engine isn't sure who you are, it won't risk naming you.

Structured Data

Structured data (schema markup) is the machine-readable layer that states plainly what your business is, does, and serves. When an engine retrieves a page to answer a question, schema lets it extract facts cleanly instead of inferring them from prose. Clean extraction means confident citation. Missing schema means the engine has to guess — and cautious engines skip guesses.

Citation Velocity

This is the modern heir to the backlink. Citation velocity measures how frequently and how recently credible third-party sources mention your business by name. Crucially, in the AI era, a mention counts even without a hyperlink. A recent surge of relevant, quality mentions signals to the engine that your business is active, relevant, and worth surfacing now.

Review Signals

AI engines read reviews as live evidence of trust. It's not only the star rating — it's volume, recency, and sentiment across multiple platforms. A business with 200 reviews averaging 4.8, refreshed monthly, reads as a healthy operation. A business with 12 reviews from three years ago reads as dormant, regardless of how good it once was.

Content Freshness

Engines favor sources that are current. Content updated regularly, dated clearly, and written to directly answer questions signals that your information can be trusted today. Stale content is a liability in a system that prizes recency.

How the Platforms Differ

The signals are shared, but each engine emphasizes them differently.

Perplexity is retrieval-first and citation-obsessed. It leans hard on live web results and shows its sources explicitly. This makes citation velocity and clean, extractable content especially powerful for Perplexity visibility. If your pages answer the question directly and you're mentioned across credible sources, Perplexity finds you fast.

ChatGPT blends its training knowledge with live browsing. Its search layer, now used by hundreds of millions of people weekly, pulls fresh results, but its baseline answers still lean on what the model learned in training. This dual nature rewards businesses with both established entity recognition and a current, well-structured web presence.

Google Gemini and AI Overviews sit on top of Google's index and knowledge graph. Here, entity authority and structured data carry outsized weight, because Google already maintains a formal understanding of entities. If Google recognizes you as an entity and your schema is clean, you are well-positioned for Gemini and Overviews. The traditional Google trust signals — E-E-A-T, technical health — still feed this layer.

The throughline: optimize the shared signals and you improve across all three at once. Chase one platform's quirks and you build fragile visibility.

What Operators Can Actually Do

Start by measuring reality. Ask each engine what it knows about your business and who it recommends in your category. Note where you're absent.

Then work the signals in order of leverage. Establish entity clarity first — consistent name, address, phone, and verified profiles everywhere. Implement structured data so engines can extract you cleanly. Build citation velocity through genuine third-party mentions: press, partnerships, directories, guest contributions. Keep reviews current across platforms. And maintain content that answers real questions and shows its freshness.

This is coordinated infrastructure work, not a one-time fix — which is exactly how we structure [our services]([/services]) at Hey Pearl.

Frequently Asked Questions

Can I pay to be recommended by ChatGPT or Perplexity?

Not in the organic answer layer. These recommendations are earned through the signals above, not bought. Some platforms are testing ads, but the recommendations users trust most are the organic ones — and those respond to authority, not budget.

Why does one engine recommend me and another doesn't?

Because they weight signals differently. If Perplexity names you but Gemini doesn't, you likely have strong live citations but weak entity recognition in Google's knowledge graph. The gap tells you where to work.

How important are reviews really?

Very. Review recency and volume are among the clearest live-trust signals an engine can read. Consistent review activity often separates two otherwise similar businesses in AI recommendations.

Do backlinks still matter?

Links still help, but the emphasis has shifted to mentions. AI engines can attribute authority to your business from an unlinked mention in a credible source. Volume and quality of mentions — linked or not — now drive citation velocity.

How often do these engines update what they know?

The live retrieval layer updates constantly, so fresh signals can influence answers within days or weeks. The training layer updates on the model provider's schedule, so deeper recognition builds over months.

Is this worth it if I only serve a small local market?

Often more so. Local queries are exactly where AI engines produce short, named recommendations — and local markets have fewer entities competing for entity authority. Clarity wins quickly at the local level.