

How to Audit Your AI Search Visibility in Under 30 Minutes

May 22, 2026 · 8 min read · Hey Pearl Agency LLC

Most business owners have never checked what AI engines say about them. Here is a repeatable 30-minute audit that tells you exactly where you stand — and what to fix first.

You are already being described by machines

By mid-2026, a majority of buyer research journeys begin with a generative answer rather than a list of blue links. Google's AI Overviews now appear on the bulk of commercial queries, and standalone assistants — ChatGPT, Perplexity, Gemini, Claude — handle a growing share of the questions your prospects used to type into a search bar. Each of these systems already has an opinion about your business, your category, and your competitors. The problem is that almost no operator has actually read that opinion.

An AI visibility audit is the fastest way to close that gap. It does not require software, a subscription, or a consultant. It requires thirty focused minutes and a willingness to look at unflattering results honestly. This is the process we run before we ever quote a client.

Before you start: set the conditions

Run the audit in a clean session. Log out of any account that personalizes results, or use a private window, so the engine is not simply reflecting your own history back at you. Turn off any memory or custom-instruction features. You want to see what a stranger sees, not what the model has learned about you specifically.

Keep a simple notepad open with four columns: the engine, the query, whether you were mentioned, and whether the information was correct. That table is your audit.

Step one: the four engines, the same questions (about 15 minutes)

Run the same set of prompts across ChatGPT, Perplexity, Gemini, and Google AI Overviews. Using identical prompts lets you compare coverage across engines rather than testing four different things.

The queries that matter

Start with three categories of prompt:

- **Category queries.** "Who are the best [your service] providers in [your city]?" and "What companies should I consider for [the problem you solve]?" These reveal whether you appear in the consideration set at all.
- **Named queries.** "Tell me about [your business name]" and "Is [your business name] reputable?" These reveal what the engines have actually ingested about you — and whether it is accurate.
- **Problem queries.** "How do I [solve the problem your customer has]?" followed by "Who can help me with this?" This is how real buyers behave: they describe a problem before they name a vendor.

Run each query, then ask one follow-up: "Why did you recommend those?" The reasoning the engine offers is often more useful than the list itself, because it exposes the signals it weighted.

Step two: read the results like an operator (about 10 minutes)

You are looking for four things.

Presence. Did you appear? In which engines? For which query types? Appearing in named queries but never in category queries is common — it means the machines know you exist but do not yet consider you a credible answer.

Accuracy. Is your service description right? Your location? Your specialty? Outdated or wrong information is worse than absence, because it erodes trust at the exact moment of consideration.

Framing. How are you described relative to competitors? Are you the safe default, a niche option, or an afterthought? Note the adjectives the engine uses. They are downstream of your reputation signals.

Citations. In Perplexity and AI Overviews especially, note which sources the engine links to. If competitors are cited from their own sites and you are cited from a third-party directory — or not at all — that is a content architecture problem you can act on.

Red flags to write down

- You do not appear in any category query, only when named directly.
- The engine confuses you with a similarly named business.
- Your founding date, location, or core service is wrong.
- Every citation about your category points to competitors' owned pages.
- The model hedges ("I don't have specific information about...") when asked about you by name.
- Your reviews or reputation are described in neutral-to-negative terms while competitors are described warmly.

Any two of these together mean you are effectively invisible in the channel where a large and growing share of buying decisions now begin.

Step three: score it (about 5 minutes)

Use a simple 20-point framework. Award up to five points in each of four dimensions:

The 4-part visibility score

1. **Presence (0–5).** Zero if you never appear; five if you appear in category, named, and problem queries across all four engines.
2. **Accuracy (0–5).** Zero if core facts are wrong; five if every engine describes you correctly.
3. **Positioning (0–5).** Zero if you are an afterthought; five if you are framed as a leading, credible answer.
4. **Citations (0–5).** Zero if no engine cites your owned properties; five if your own pages are consistently the source.

Add the columns. **16–20** means you are ahead of your market — protect the position. **9–15** means you have a foothold but leak authority to competitors — you have clear, winnable work. **0–8** means you are functionally absent, and every week you wait is compounding advantage for whoever moves first.

Quick wins you can start this week

If your score exposed gaps, the highest-leverage fixes are almost always structural, not cosmetic:

- **Fix the facts first.** Correct wrong information on the sources engines cite — your site, your Google Business Profile, and the top directories. Accuracy is the cheapest point to buy back.
- **Publish the answer to your top problem query** as a clear, well-structured page on your own domain. Engines cite structured, specific, first-party content.
- **Establish a clean entity.** Consistent name, address, description, and founding details across every property you control. Machines reward coherence.
- **Address reputation signals.** Sentiment in reviews increasingly feeds how engines frame you. A deliberate review strategy moves positioning over a quarter.

Re-run the audit every 90 days. AI visibility is not a project you finish; it is a position you hold. The operators who treat it as a standing discipline — the way they already treat cash flow or hiring — are the ones the machines will keep recommending.

You cannot improve what you have never looked at. Thirty minutes of honest auditing tells you more about your position in the AI search era than most businesses will learn in a year of guessing. Run the audit, score it, fix the highest-leverage gaps, and re-run it next quarter. If the results surprised you — or if you want the structural work done properly rather than patched — this is exactly the discipline we

build for our clients.

Frequently Asked Questions

How often should I run an AI visibility audit?

Every 90 days is a sensible cadence for most businesses. AI engines update their models and re-index sources continuously, and your competitors are moving. A quarterly audit lets you track whether your score is climbing, holding, or slipping, and it aligns with the roughly one-quarter horizon it takes structural fixes to register.

Do I need paid tools to audit my AI visibility?

No. The core audit requires only free access to ChatGPT, Perplexity, Gemini, and Google AI Overviews, plus a notepad to record results. Paid monitoring platforms are useful for tracking at scale over time, but the manual 30-minute audit is enough to establish where you stand and what to fix first.

Why should I run queries while logged out?

Personalization and memory features cause the engines to reflect your own history back at you, which produces flattering but misleading results. A logged-out or private session shows you what a prospect who has never heard of you actually sees, which is the only view that matters for acquisition.

What if the AI engines have my information wrong?

Inaccurate information is a priority fix and often the cheapest point to recover. Correct the facts on the sources these engines cite most — your own website, your Google Business Profile, and major directories. Ensure your entity details are consistent everywhere, then re-audit in a few weeks to confirm the correction has propagated.

Why do I appear when named but not in category searches?

It means the engines know you exist but do not yet consider you a credible answer to the buyer's underlying problem. Closing that gap requires authority signals: first-party content that answers category and problem queries, consistent entity data, and reputation signals that position you as a leading option rather than a footnote.

Which engine matters most for my business?

It depends on your buyers, but you should not optimize for one. Google AI Overviews reach the broadest audience, Perplexity is favored by research-heavy buyers, and ChatGPT and Gemini carry enormous conversational query volume. The good news is that the underlying signals — accurate entities, structured first-party content, strong reputation — improve visibility across all of them at once.

What is a good AI visibility score?

On the 20-point framework, a score of 16 or above means you are ahead of your market and should focus on defending the position. Nine to 15 indicates a foothold with clear, winnable gaps. Eight or below means you are effectively invisible in AI search and should treat it as an urgent priority, since first movers in a category compound their advantage.