

llms.txt and AI Crawler Access: How to Control What AI Engines Read on Your Site

Published 2026-08-20 · 10 min read

Before AI engines can cite your business, their crawlers have to be able to read it. Here is what llms.txt is, which AI bots are already visiting your site, and how to configure crawler access so you are found instead of filtered out.

Why AI crawler access is the step most businesses skip

Every AI visibility strategy — schema markup, entity architecture, authority content, reputation signals — depends on a single precondition: the AI engine has to be able to crawl and read your site in the first place.

Most businesses invest heavily in the content and structured data layer while never checking whether the underlying crawler permissions actually allow AI systems to reach that content. It is a foundational gap, and it is invisible until you specifically look for it. Your website can be technically flawless and your content genuinely citable, and still be functionally invisible to ChatGPT, Perplexity, or Google AI Overviews if the crawlers behind them are blocked, throttled, or simply never told the site exists.

This is the layer beneath the layer — the access question that has to be answered correctly before any other AI visibility work can compound.

What llms.txt actually is

llms.txt is a proposed plain-text file, placed at a website's root (yourdomain.com/llms.txt), designed to give large language models a curated, structured overview of a site's most important content — written in a format optimized for machine consumption rather than human browsing.

Conceptually, it sits between a sitemap and a README file. Where a sitemap.xml lists every URL for traditional search engine indexing, llms.txt is meant to guide an AI system directly to the pages that matter most: your core service pages, your most authoritative articles, your FAQ content, your about and contact information. It typically uses simple Markdown formatting — a title, a short description, and organized links grouped by section — so a language model can quickly orient itself around what a site is and where its most citable content lives.

It is important to be precise about its status: llms.txt is a community-proposed convention, not a standard formally adopted or confirmed to be used by OpenAI, Anthropic, Google, or Perplexity in the way robots.txt is universally respected by search crawlers. Adoption is inconsistent and evolving. That does not make it worthless — a well-structured llms.txt costs little to implement, causes no harm, and positions a site to benefit if and as adoption grows. But it should be treated as a forward-looking addition, not a substitute for the parts of AI crawler access that are already confirmed and consequential: robots.txt directives and server-level access rules.

llms.txt vs. robots.txt vs. sitemap.xml — what each one actually does

These three files are often confused because they all live at the root of a domain, but they serve distinct functions:

robots.txt is the access-control layer. It tells crawlers — search engines and AI bots alike — which parts of your site they are allowed or forbidden to request. This is the file with real, immediate teeth: a disallow directive aimed at an AI crawler will, in most cases, be honored, meaning that crawler simply will not fetch the blocked pages.

sitemap.xml is the discovery layer for traditional search engines. It lists every indexable URL so Google, Bing, and other search crawlers can find and prioritize content efficiently. It is not primarily designed with AI language model consumption in mind, though AI systems that also function as search engines (like Google's AI Overviews) benefit from it indirectly.

llms.txt is the orientation layer, specifically for AI systems. It does not control access — a bot can ignore it entirely and suffer no consequence — but where it is respected, it shortcuts an AI system's understanding of what your site is and which pages represent your most authoritative, most current content.

The practical takeaway: robots.txt determines whether AI crawlers can reach your content at all. llms.txt, where honored, determines how efficiently they understand it once they arrive. Get the access layer right first.

The AI crawlers you need to know by name

Every major AI company operates named, identifiable web crawlers. Configuring access requires knowing which bot serves which purpose:

GPTBot — OpenAI's crawler, used to gather data for training and improving its models. A separate crawler, OAI-SearchBot, handles retrieval for ChatGPT's real-time search and browsing features, and is generally the one you want indexing your content for citation purposes.

ChatGPT-User — A distinct user-agent OpenAI uses when ChatGPT is actively browsing a page on behalf of a live user query, as opposed to bulk training-data collection.

ClaudeBot — Anthropic's crawler, used for both training data collection and to support Claude's ability to browse and reference web content when responding to queries.

PerplexityBot — Perplexity's crawler, directly responsible for the real-time retrieval that powers Perplexity's answer engine. Because Perplexity operates primarily as a citation-driven answer engine rather than a pure chatbot, this crawler has an outsized effect on whether your business appears as a cited source.

Google-Extended — A control token (not a traditional crawler user-agent) that lets you specifically opt your content in or out of use for training Google's Gemini models and powering AI Overviews, separately from standard Googlebot indexing for search.

Amazonbot — Amazon's crawler, which feeds Alexa's answer capabilities among other uses.

Applebot-Extended — Apple's control token for AI-related use of crawled content, relevant as Apple integrates more generative AI features into its ecosystem.

New crawlers are introduced as AI companies expand capabilities, so this list should be treated as a living reference, not a permanent one. The businesses that stay current are the ones who periodically re-check which

bots are actually requesting their pages.

How to check whether AI crawlers can currently access your site

Before writing a single directive, establish your current baseline:

Check your existing robots.txt. Visit `yourdomain.com/robots.txt` directly. Read every line. Many robots.txt files were written years before AI crawlers existed and contain broad disallow rules — intended for a different purpose entirely — that inadvertently catch AI bots in their net.

Check your server logs. If you have access to raw server or CDN logs, search for the crawler names above. Their presence confirms they are attempting to reach your site; their absence tells you nothing has been actively blocking them, but also that nothing has actively invited them either.

Check your CDN or firewall rules. Services like Cloudflare have introduced bot-management features that can block AI crawlers by default at the network level — sometimes without a business realizing the setting exists or was ever changed. This layer sits outside robots.txt entirely and can silently override an otherwise-permissive robots.txt file.

Use a live fetch test. Some AI platforms offer or are rumored to offer tools that let you test whether their crawler can successfully fetch a given URL. Where available, use them directly rather than assuming.

How to write an llms.txt file

A basic llms.txt follows a simple, readable Markdown structure:

```
# [Your Business Name]

> One or two sentence description of what the business does and who it serves.

## Core Pages
- [About](https://yourdomain.com/about): Company background and team
- [Services](https://yourdomain.com/services): Full service overview

## Authority Content
- [Guide Title](https://yourdomain.com/insights/guide-slug): One-line summary
- [FAQ](https://yourdomain.com/faq): Answers to common questions

## Contact
- [Contact](https://yourdomain.com/contact): How to reach the business
```

Keep descriptions factual and specific rather than promotional — the purpose is orientation, not marketing copy. Prioritize your highest-authority, most frequently updated content over comprehensive completeness; a curated list of twenty genuinely important pages is more useful to a language model than an exhaustive dump of every URL on the site.

How to configure robots.txt for AI crawlers

The strategic decision is not simply "allow" or "block" — it depends on what each crawler is used for and what your visibility goals are.

For businesses pursuing AI visibility and citation (the default recommendation for nearly every commercial business), the general posture should be permissive toward retrieval-oriented crawlers:

```
User-agent: OAI-SearchBot
Allow: /

User-agent: PerplexityBot
Allow: /

User-agent: ClaudeBot
Allow: /

User-agent: Google-Extended
Allow: /
```

For businesses with specific content they do not want used in AI training — while still wanting to appear in AI-generated answers — the distinction between training crawlers and retrieval crawlers matters. You can, for example, disallow GPTBot (training) while allowing OAI-SearchBot (real-time retrieval and citation), giving you citation visibility without contributing to model training corpora.

Directories that should generally remain disallowed for all crawlers regardless of AI strategy: admin panels, staging environments, internal search result pages, checkout and account pages, and any duplicate or thin-content page groups you would not want cited as a representative sample of your business.

Always test syntax after editing — a single misplaced character in robots.txt can invalidate the entire file or accidentally disallow everything.

Common mistakes that quietly block AI visibility

Inherited disallow-all rules. Some CMS platforms and site builders ship with a default `Disallow: /` in staging or preview environments that gets carried into production and never removed.

Blanket bot-blocking plugins. Security plugins marketed as blocking "bad bots" or "scrapers" sometimes categorize all automated crawlers — including legitimate AI retrieval bots — as threats by default, with no distinction made.

CDN-level blocking that bypasses robots.txt entirely. As noted above, this is the mistake most likely to go unnoticed, because checking robots.txt alone will show everything looks fine while a network-level rule silently drops the requests before they ever reach your server.

Treating llms.txt as sufficient on its own. A perfectly written llms.txt does nothing if robots.txt or a firewall rule is blocking the crawler from ever requesting the site in the first place. Access has to be granted before orientation matters.

Does blocking AI crawlers protect your content, or hurt your visibility?

This is a legitimate strategic question, not just a technical one, and the honest answer is: it depends on the goal.

Blocking AI training crawlers (GPTBot, ClaudeBot's training use, Google-Extended) does prevent your specific content from being used as raw material in future model training runs. For businesses with significant proprietary

or licensed content, that may be a deliberate and reasonable choice.

However, blocking retrieval and citation crawlers (OAI-SearchBot, PerplexityBot's live retrieval function) has a different effect entirely: it prevents your business from appearing in the AI-generated answers and recommendations those systems produce in response to real user queries. For a business trying to be the one AI recommends, blocking these crawlers is directly counterproductive — it removes you from consideration before your content and authority signals ever get evaluated.

The strategic default for most commercial businesses: allow retrieval and citation crawlers broadly, and make training-data decisions deliberately and separately, rather than blocking everything under a single blunt "block all AI bots" instinct.

Where this fits in the larger AI authority system

Crawler access is not a standalone tactic — it is the infrastructure layer everything else depends on. Entity architecture, schema markup, and authority content all assume the AI engine can actually reach and parse the page carrying those signals. A business with excellent structured data and a misconfigured robots.txt file is investing in signals that never get received.

This is why crawler access belongs at the start of an AI visibility audit, not as an afterthought. Before evaluating content quality or entity completeness, confirm the door is actually open.

Frequently Asked Questions

What is llms.txt used for?

llms.txt is a plain-text file placed at a website's root domain that gives AI language models a curated, structured summary of a site's most important pages — its core services, authority content, and contact information — in a format designed for quick machine orientation rather than human browsing. It functions similarly to a sitemap, but is written specifically to help AI systems understand what a site is and where its best content lives.

Is llms.txt an official web standard?

No. It is a community-proposed convention, not a standard formally confirmed to be used by OpenAI, Anthropic, Google, or Perplexity the way robots.txt is universally respected by search crawlers. Adoption is inconsistent and still evolving. It is worth implementing because it costs little and positions a site to benefit as adoption grows, but it should not be treated as a substitute for the access controls in robots.txt, which are confirmed and consequential.

Should I block AI crawlers like GPTBot from scraping my content?

It depends on your goal. Blocking training-oriented crawlers (like GPTBot's training use) prevents your content from being used in future model training, which may be reasonable for businesses with proprietary content. But blocking retrieval and citation crawlers, such as OAI-SearchBot or PerplexityBot, prevents your business from appearing in AI-generated answers and recommendations — which is directly counterproductive for a business trying to build AI visibility. Most commercial businesses should allow retrieval crawlers broadly and make training-data decisions separately and deliberately.

How do I know which AI bots are currently crawling my website?

Check your raw server or CDN access logs for known AI crawler user-agent names — GPTBot, OAI-SearchBot, ChatGPT-User, ClaudeBot, PerplexityBot, Amazonbot, and similar. Their presence in logs confirms they are actively requesting pages. You should also review your robots.txt file directly and check whether your CDN or firewall (such as Cloudflare) has bot-management rules that could be blocking crawlers at the network level, independent of what robots.txt says.

Does llms.txt replace robots.txt?

No — they serve different functions and both matter. robots.txt is the access-control layer; it determines whether a crawler is permitted to request your pages at all, and is treated as authoritative by essentially every major crawler. llms.txt is an orientation layer that, where honored, helps an AI system understand your site's structure and priority content more efficiently once access has already been granted. Access has to be confirmed first — an llms.txt file does nothing for a crawler that robots.txt or a firewall rule is blocking outright.

Will blocking AI crawlers hurt my AI search visibility?

If you block the retrieval and citation crawlers that power live AI answers — such as PerplexityBot or OAI-SearchBot — yes, directly. Those crawlers are how AI engines discover and later cite your content in response to real user queries. A business blocking these crawlers is removing itself from consideration in AI-generated recommendations before any content or authority signal is ever evaluated, regardless of how strong that content otherwise is.

Conclusion

AI crawler access is the layer every other AI visibility tactic depends on, and it is the one most audits skip. A business can have excellent schema, strong entity signals, and genuinely citable content, and still be invisible to AI engines if a stale robots.txt rule, a default security plugin, or a CDN-level bot filter is quietly blocking the crawlers that would otherwise find it. Confirming access, writing a clear llms.txt, and making deliberate — not accidental — decisions about which AI bots can reach your site is foundational infrastructure work. It does not replace content or entity strategy, but without it, that other work never gets the chance to be seen.